

SL Project - Fine-tuning

December 14, 2022

```
[1]: import torchvision
import torch
import torch.nn as nn
from torchsummary import summary
```

```
[2]: from torchvision import transforms, datasets
import matplotlib.pyplot as plt
from tqdm import tqdm
```

```
[3]: import pandas as pd
```

```
[4]: import os
import pickle
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[5]: sns.set_style("darkgrid")
```

```
[6]: import cv2

image = cv2.imread("../data/asl_alphabet_train/A/A1.jpg")
print(image.shape)  # image dimensions

(200, 200, 3)
```

```
[7]: transform = transforms.Compose([
    transforms.Resize(256),
    transforms.RandomCrop(224),
    transforms.ToTensor()
])
```

```
[8]: PATH = "../data/asl_alphabet_train/"
```

```
[9]: dataset = datasets.ImageFolder(PATH, transform=transform)
```

```
[10]: n = len(dataset)
```

```
[11]: print(n)
```

87000

```
[12]: torch.manual_seed(1)
      indices = torch.randperm(n)
```

```
[13]: test_proportion = 0.2 # 20 percent of data used for testing
      test_size = int(n * test_proportion)
```

```
[14]: train_dataset = torch.utils.data.Subset(dataset, indices[test_size:])
      test_dataset = torch.utils.data.Subset(dataset, indices[:test_size])
```

```
[15]: len(train_dataset)
```

69600

```
[16]: len(test_dataset)
```

17400

```
[17]: train_dataloader = torch.utils.data.DataLoader(dataset=train_dataset,
                                                    batch_size=32,
                                                    shuffle=True,
                                                    num_workers=4)

      test_dataloader = torch.utils.data.DataLoader(dataset=test_dataset,
                                                    batch_size=32,
                                                    shuffle=False,
                                                    num_workers=4)
```

```
[18]: classes = dataset.classes
```

```
[19]: cols = 10
      rows = 2
      fig, ax = plt.subplots(rows, cols, figsize=(20, 4))
      i = 0
      for img, label in train_dataloader:
          plt.subplot(rows, cols, i + 1)
          plt.imshow(img[0].permute(1, 2, 0))
          plt.xticks(())
          plt.yticks(())
          plt.title(classes[label[0]])
          i += 1
      if i == 20:
          break
```



```
[20]: model = torch.hub.load(repo_or_dir='pytorch/vision:v0.10.0', model='googlenet',
                             weights='GoogLeNet_Weights.IMAGENET1K_V1')
```

Using cache found in /home/rao.ans/.cache/torch/hub/pytorch_vision_v0.10.0

```
[21]: model.fc = torch.nn.Linear(model.fc.in_features, len(classes))
```

```
[22]: criterion = torch.nn.CrossEntropyLoss()
```

```
[23]: optimizer = torch.optim.Adam(model.parameters(), lr=3e-4, weight_decay=0.001)
```

```
[24]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
[25]: if torch.cuda.is_available():
        model.cuda()
```

```
[26]: summary(model, (3, 224, 224))
```

Layer (type)	Output Shape	Param #

Conv2d-1	[-1, 64, 112, 112]	9,408
BatchNorm2d-2	[-1, 64, 112, 112]	128
BasicConv2d-3	[-1, 64, 112, 112]	0
MaxPool2d-4	[-1, 64, 56, 56]	0
Conv2d-5	[-1, 64, 56, 56]	4,096
BatchNorm2d-6	[-1, 64, 56, 56]	128
BasicConv2d-7	[-1, 64, 56, 56]	0
Conv2d-8	[-1, 192, 56, 56]	110,592
BatchNorm2d-9	[-1, 192, 56, 56]	384
BasicConv2d-10	[-1, 192, 56, 56]	0
MaxPool2d-11	[-1, 192, 28, 28]	0
Conv2d-12	[-1, 64, 28, 28]	12,288
BatchNorm2d-13	[-1, 64, 28, 28]	128
BasicConv2d-14	[-1, 64, 28, 28]	0
Conv2d-15	[-1, 96, 28, 28]	18,432
BatchNorm2d-16	[-1, 96, 28, 28]	192
BasicConv2d-17	[-1, 96, 28, 28]	0

Conv2d-18	[-1, 128, 28, 28]	110,592
BatchNorm2d-19	[-1, 128, 28, 28]	256
BasicConv2d-20	[-1, 128, 28, 28]	0
Conv2d-21	[-1, 16, 28, 28]	3,072
BatchNorm2d-22	[-1, 16, 28, 28]	32
BasicConv2d-23	[-1, 16, 28, 28]	0
Conv2d-24	[-1, 32, 28, 28]	4,608
BatchNorm2d-25	[-1, 32, 28, 28]	64
BasicConv2d-26	[-1, 32, 28, 28]	0
MaxPool2d-27	[-1, 192, 28, 28]	0
Conv2d-28	[-1, 32, 28, 28]	6,144
BatchNorm2d-29	[-1, 32, 28, 28]	64
BasicConv2d-30	[-1, 32, 28, 28]	0
Inception-31	[-1, 256, 28, 28]	0
Conv2d-32	[-1, 128, 28, 28]	32,768
BatchNorm2d-33	[-1, 128, 28, 28]	256
BasicConv2d-34	[-1, 128, 28, 28]	0
Conv2d-35	[-1, 128, 28, 28]	32,768
BatchNorm2d-36	[-1, 128, 28, 28]	256
BasicConv2d-37	[-1, 128, 28, 28]	0
Conv2d-38	[-1, 192, 28, 28]	221,184
BatchNorm2d-39	[-1, 192, 28, 28]	384
BasicConv2d-40	[-1, 192, 28, 28]	0
Conv2d-41	[-1, 32, 28, 28]	8,192
BatchNorm2d-42	[-1, 32, 28, 28]	64
BasicConv2d-43	[-1, 32, 28, 28]	0
Conv2d-44	[-1, 96, 28, 28]	27,648
BatchNorm2d-45	[-1, 96, 28, 28]	192
BasicConv2d-46	[-1, 96, 28, 28]	0
MaxPool2d-47	[-1, 256, 28, 28]	0
Conv2d-48	[-1, 64, 28, 28]	16,384
BatchNorm2d-49	[-1, 64, 28, 28]	128
BasicConv2d-50	[-1, 64, 28, 28]	0
Inception-51	[-1, 480, 28, 28]	0
MaxPool2d-52	[-1, 480, 14, 14]	0
Conv2d-53	[-1, 192, 14, 14]	92,160
BatchNorm2d-54	[-1, 192, 14, 14]	384
BasicConv2d-55	[-1, 192, 14, 14]	0
Conv2d-56	[-1, 96, 14, 14]	46,080
BatchNorm2d-57	[-1, 96, 14, 14]	192
BasicConv2d-58	[-1, 96, 14, 14]	0
Conv2d-59	[-1, 208, 14, 14]	179,712
BatchNorm2d-60	[-1, 208, 14, 14]	416
BasicConv2d-61	[-1, 208, 14, 14]	0
Conv2d-62	[-1, 16, 14, 14]	7,680
BatchNorm2d-63	[-1, 16, 14, 14]	32
BasicConv2d-64	[-1, 16, 14, 14]	0
Conv2d-65	[-1, 48, 14, 14]	6,912

BatchNorm2d-66	[-1, 48, 14, 14]	96
BasicConv2d-67	[-1, 48, 14, 14]	0
MaxPool2d-68	[-1, 480, 14, 14]	0
Conv2d-69	[-1, 64, 14, 14]	30,720
BatchNorm2d-70	[-1, 64, 14, 14]	128
BasicConv2d-71	[-1, 64, 14, 14]	0
Inception-72	[-1, 512, 14, 14]	0
Conv2d-73	[-1, 160, 14, 14]	81,920
BatchNorm2d-74	[-1, 160, 14, 14]	320
BasicConv2d-75	[-1, 160, 14, 14]	0
Conv2d-76	[-1, 112, 14, 14]	57,344
BatchNorm2d-77	[-1, 112, 14, 14]	224
BasicConv2d-78	[-1, 112, 14, 14]	0
Conv2d-79	[-1, 224, 14, 14]	225,792
BatchNorm2d-80	[-1, 224, 14, 14]	448
BasicConv2d-81	[-1, 224, 14, 14]	0
Conv2d-82	[-1, 24, 14, 14]	12,288
BatchNorm2d-83	[-1, 24, 14, 14]	48
BasicConv2d-84	[-1, 24, 14, 14]	0
Conv2d-85	[-1, 64, 14, 14]	13,824
BatchNorm2d-86	[-1, 64, 14, 14]	128
BasicConv2d-87	[-1, 64, 14, 14]	0
MaxPool2d-88	[-1, 512, 14, 14]	0
Conv2d-89	[-1, 64, 14, 14]	32,768
BatchNorm2d-90	[-1, 64, 14, 14]	128
BasicConv2d-91	[-1, 64, 14, 14]	0
Inception-92	[-1, 512, 14, 14]	0
Conv2d-93	[-1, 128, 14, 14]	65,536
BatchNorm2d-94	[-1, 128, 14, 14]	256
BasicConv2d-95	[-1, 128, 14, 14]	0
Conv2d-96	[-1, 128, 14, 14]	65,536
BatchNorm2d-97	[-1, 128, 14, 14]	256
BasicConv2d-98	[-1, 128, 14, 14]	0
Conv2d-99	[-1, 256, 14, 14]	294,912
BatchNorm2d-100	[-1, 256, 14, 14]	512
BasicConv2d-101	[-1, 256, 14, 14]	0
Conv2d-102	[-1, 24, 14, 14]	12,288
BatchNorm2d-103	[-1, 24, 14, 14]	48
BasicConv2d-104	[-1, 24, 14, 14]	0
Conv2d-105	[-1, 64, 14, 14]	13,824
BatchNorm2d-106	[-1, 64, 14, 14]	128
BasicConv2d-107	[-1, 64, 14, 14]	0
MaxPool2d-108	[-1, 512, 14, 14]	0
Conv2d-109	[-1, 64, 14, 14]	32,768
BatchNorm2d-110	[-1, 64, 14, 14]	128
BasicConv2d-111	[-1, 64, 14, 14]	0
Inception-112	[-1, 512, 14, 14]	0
Conv2d-113	[-1, 112, 14, 14]	57,344

BatchNorm2d-114	[-1, 112, 14, 14]	224
BasicConv2d-115	[-1, 112, 14, 14]	0
Conv2d-116	[-1, 144, 14, 14]	73,728
BatchNorm2d-117	[-1, 144, 14, 14]	288
BasicConv2d-118	[-1, 144, 14, 14]	0
Conv2d-119	[-1, 288, 14, 14]	373,248
BatchNorm2d-120	[-1, 288, 14, 14]	576
BasicConv2d-121	[-1, 288, 14, 14]	0
Conv2d-122	[-1, 32, 14, 14]	16,384
BatchNorm2d-123	[-1, 32, 14, 14]	64
BasicConv2d-124	[-1, 32, 14, 14]	0
Conv2d-125	[-1, 64, 14, 14]	18,432
BatchNorm2d-126	[-1, 64, 14, 14]	128
BasicConv2d-127	[-1, 64, 14, 14]	0
MaxPool2d-128	[-1, 512, 14, 14]	0
Conv2d-129	[-1, 64, 14, 14]	32,768
BatchNorm2d-130	[-1, 64, 14, 14]	128
BasicConv2d-131	[-1, 64, 14, 14]	0
Inception-132	[-1, 528, 14, 14]	0
Conv2d-133	[-1, 256, 14, 14]	135,168
BatchNorm2d-134	[-1, 256, 14, 14]	512
BasicConv2d-135	[-1, 256, 14, 14]	0
Conv2d-136	[-1, 160, 14, 14]	84,480
BatchNorm2d-137	[-1, 160, 14, 14]	320
BasicConv2d-138	[-1, 160, 14, 14]	0
Conv2d-139	[-1, 320, 14, 14]	460,800
BatchNorm2d-140	[-1, 320, 14, 14]	640
BasicConv2d-141	[-1, 320, 14, 14]	0
Conv2d-142	[-1, 32, 14, 14]	16,896
BatchNorm2d-143	[-1, 32, 14, 14]	64
BasicConv2d-144	[-1, 32, 14, 14]	0
Conv2d-145	[-1, 128, 14, 14]	36,864
BatchNorm2d-146	[-1, 128, 14, 14]	256
BasicConv2d-147	[-1, 128, 14, 14]	0
MaxPool2d-148	[-1, 528, 14, 14]	0
Conv2d-149	[-1, 128, 14, 14]	67,584
BatchNorm2d-150	[-1, 128, 14, 14]	256
BasicConv2d-151	[-1, 128, 14, 14]	0
Inception-152	[-1, 832, 14, 14]	0
MaxPool2d-153	[-1, 832, 7, 7]	0
Conv2d-154	[-1, 256, 7, 7]	212,992
BatchNorm2d-155	[-1, 256, 7, 7]	512
BasicConv2d-156	[-1, 256, 7, 7]	0
Conv2d-157	[-1, 160, 7, 7]	133,120
BatchNorm2d-158	[-1, 160, 7, 7]	320
BasicConv2d-159	[-1, 160, 7, 7]	0
Conv2d-160	[-1, 320, 7, 7]	460,800
BatchNorm2d-161	[-1, 320, 7, 7]	640

BasicConv2d-162	[-1, 320, 7, 7]	0
Conv2d-163	[-1, 32, 7, 7]	26,624
BatchNorm2d-164	[-1, 32, 7, 7]	64
BasicConv2d-165	[-1, 32, 7, 7]	0
Conv2d-166	[-1, 128, 7, 7]	36,864
BatchNorm2d-167	[-1, 128, 7, 7]	256
BasicConv2d-168	[-1, 128, 7, 7]	0
MaxPool2d-169	[-1, 832, 7, 7]	0
Conv2d-170	[-1, 128, 7, 7]	106,496
BatchNorm2d-171	[-1, 128, 7, 7]	256
BasicConv2d-172	[-1, 128, 7, 7]	0
Inception-173	[-1, 832, 7, 7]	0
Conv2d-174	[-1, 384, 7, 7]	319,488
BatchNorm2d-175	[-1, 384, 7, 7]	768
BasicConv2d-176	[-1, 384, 7, 7]	0
Conv2d-177	[-1, 192, 7, 7]	159,744
BatchNorm2d-178	[-1, 192, 7, 7]	384
BasicConv2d-179	[-1, 192, 7, 7]	0
Conv2d-180	[-1, 384, 7, 7]	663,552
BatchNorm2d-181	[-1, 384, 7, 7]	768
BasicConv2d-182	[-1, 384, 7, 7]	0
Conv2d-183	[-1, 48, 7, 7]	39,936
BatchNorm2d-184	[-1, 48, 7, 7]	96
BasicConv2d-185	[-1, 48, 7, 7]	0
Conv2d-186	[-1, 128, 7, 7]	55,296
BatchNorm2d-187	[-1, 128, 7, 7]	256
BasicConv2d-188	[-1, 128, 7, 7]	0
MaxPool2d-189	[-1, 832, 7, 7]	0
Conv2d-190	[-1, 128, 7, 7]	106,496
BatchNorm2d-191	[-1, 128, 7, 7]	256
BasicConv2d-192	[-1, 128, 7, 7]	0
Inception-193	[-1, 1024, 7, 7]	0
AdaptiveAvgPool2d-194	[-1, 1024, 1, 1]	0
Dropout-195	[-1, 1024]	0
Linear-196	[-1, 29]	29,725

=====
Total params: 5,629,629

Trainable params: 5,629,629

Non-trainable params: 0

Input size (MB): 0.57

Forward/backward pass size (MB): 94.10

Params size (MB): 21.48

Estimated Total Size (MB): 116.15

```
[27]: train_losses = []
      train_accuracies = []
      val_losses = []
      val_accuracies = []
```

```
[28]: step = 0
      no_of_epochs = 10

      for epoch in tqdm(range(no_of_epochs)):
          correct_train, total_train = 0, 0
          train_loss = 0
          model.train()
          for i, (images, labels) in enumerate(train_dataloader):
              step += 1
              images = images.to(device)
              labels = labels.to(device)
              optimizer.zero_grad()
              outputs = model(images)
              loss = criterion(outputs, labels)
              loss.backward()
              optimizer.step()

              total_train += labels.size(0)
              _, predicted = torch.max(outputs, dim=1)
              correct_train += (predicted == labels).sum().item()

          if step % 1000 == 0:
              print(f"epoch [{epoch + 1}]/[{no_of_epochs}]", end=" ")
              train_loss += loss.item()
              train_accuracy = (correct_train / total_train) * 100
              print(f"train accuracy: {train_accuracy}", end=" ")

              with torch.no_grad():
                  correct_val, total_val = 0, 0
                  val_loss = 0
                  model.eval()
                  for images, labels in test_dataloader:
                      images = images.to(device)
                      labels = labels.to(device)
                      outputs = model(images)

                      total_val += labels.size(0)
                      _, predicted = torch.max(outputs, dim=1)
                      correct_val += (predicted == labels).sum().item()

                  val_loss += loss.item()
```



```

val_accuracy = (correct_val / total_val) * 100
print(f"val accuracy: {val_accuracy}")
train_losses.append(train_loss / total_train)
val_losses.append(val_loss / total_val)
train_accuracies.append(train_accuracy)
val_accuracies.append(val_accuracy)

```

```

0%|          | 0/10 [00:00<?, ?it/s]

epoch [1]/[10] train accuracy: 96.50937499999999 val accuracy: 99.0919540229885
epoch [1]/[10] train accuracy: 96.8703125 val accuracy: 99.56896551724138

10%|         | 1/10 [04:25<39:50, 265.63s/it]

epoch [2]/[10] train accuracy: 99.84469696969697 val accuracy: 99.95402298850574
epoch [2]/[10] train accuracy: 98.59246575342466 val accuracy: 99.45977011494253

20%|         | 2/10 [08:55<35:44, 268.10s/it]

epoch [3]/[10] train accuracy: 99.79326923076923 val accuracy: 99.7816091954023
epoch [3]/[10] train accuracy: 98.9659090909091 val accuracy: 99.50574712643679

30%|         | 3/10 [13:25<31:21, 268.81s/it]

epoch [4]/[10] train accuracy: 99.84868421052632 val accuracy: 99.9080459770115
epoch [4]/[10] train accuracy: 99.00423728813558 val accuracy: 98.9655172413793

40%|         | 4/10 [17:43<26:27, 264.66s/it]

epoch [5]/[10] train accuracy: 99.78125 val accuracy: 99.89655172413792
epoch [5]/[10] train accuracy: 99.0 val accuracy: 99.2816091954023

50%|         | 5/10 [22:01<21:51, 262.33s/it]

epoch [6]/[10] train accuracy: 99.825 val accuracy: 99.91954022988506
epoch [6]/[10] train accuracy: 99.13888888888889 val accuracy: 99.4080459770115
epoch [6]/[10] train accuracy: 99.10441176470587 val accuracy: 99.22413793103448

60%|         | 6/10 [26:51<18:07, 271.83s/it]

epoch [7]/[10] train accuracy: 99.8092105263158 val accuracy: 99.74712643678161
epoch [7]/[10] train accuracy: 99.2451923076923 val accuracy: 99.78735632183908

70%|         | 7/10 [31:13<13:25, 268.42s/it]

epoch [8]/[10] train accuracy: 99.80645161290323 val accuracy: 99.93103448275862
epoch [8]/[10] train accuracy: 99.30281690140845 val accuracy: 98.94252873563218

80%|         | 8/10 [35:33<08:51, 265.90s/it]

epoch [9]/[10] train accuracy: 99.84895833333334 val accuracy: 99.94827586206897
epoch [9]/[10] train accuracy: 99.35546875 val accuracy: 99.54597701149426

90%|         | 9/10 [39:56<04:24, 264.97s/it]

```

```
epoch [10]/[10] train accuracy: 99.83088235294117 val accuracy:
99.94252873563218
epoch [10]/[10] train accuracy: 99.3048245614035 val accuracy: 98.34482758620689
100%|      | 10/10 [44:20<00:00, 266.03s/it]
```

```
[29]: SAVE_PATH = "../data/googlenet_asl_v2.pth"
```

```
[30]: torch.save(model, SAVE_PATH)
```


```
[31]: with open('../data/googlenet_asl_v2_train_losses.pkl', 'wb') as f:
      pickle.dump(train_losses, f)
```

```
[32]: with open('../data/googlenet_asl_v2_val_losses.pkl', 'wb') as f:
      pickle.dump(val_losses, f)
```

```
[33]: with open('../data/googlenet_asl_v2_train_accuracies.pkl', 'wb') as f:
      pickle.dump(train_accuracies, f)
```

```
[34]: with open('../data/googlenet_asl_v2_val_accuracies.pkl', 'wb') as f:
      pickle.dump(val_accuracies, f)
```

```
[35]: loss = pd.DataFrame({'train': train_losses, 'val': val_losses})
      sns.lineplot(loss)
      plt.xlabel("step")
      plt.ylabel("loss");
```



```
[36]: accuracy = pd.DataFrame({'train': train_accuracies, 'val': val_accuracies})  
sns.lineplot(accuracy)  
plt.xlabel("step")  
plt.ylabel("accuracy");
```

